

LDP two-sample chi-squared test

1 Setting

- $\mathbf{Y}_i \stackrel{iid}{\sim} \text{multi}(n_1, \mathbf{p}_Y), \mathbf{Z}_i \stackrel{iid}{\sim} \text{multi}(n_2, \mathbf{p}_Z)$ with k categories
- One-hot vector form i.e. random vectors with dependent Bernoulli random variable entries
- Allow for $n_1 \neq n_2$

2 Generalized Randomized Response and two sample Pearson chi-square statistic

2.1 Privacy mechanism: Generalized Randomized Response

Definition 2.1 (Generalized Randomized Response (Theorem 5.4. of Gaboardi and Rogers [1])).
For a multinomial random vector $\mathbf{Y}_i \stackrel{iid}{\sim} \text{multi}(n_1, \mathbf{p}_Y)$, we define

$$\mathbb{P}(\mathcal{M}_{\text{GenRR}}(\mathbf{Y}_i) = \mathbf{y}' | \mathbf{Y}_i = \mathbf{y}) := \begin{cases} \frac{\exp(\alpha)}{\exp(\alpha) + k - 1} & \text{if } \mathbf{y}' = \mathbf{y} \\ \frac{1}{\exp(\alpha) + k - 1} & \text{if } \mathbf{y}' \neq \mathbf{y}. \end{cases}$$

Then $\tilde{\mathbf{Y}}_i := \mathcal{M}_{\text{GenRR}}(\mathbf{Y}_i)$ is a multinomial random vector with probability vector

$$\tilde{\mathbf{p}}_Y := \mathbf{p}_Y \frac{\exp(\alpha)}{\exp(\alpha) + k - 1} + (1 - \mathbf{p}_Y) \frac{1}{\exp(\alpha) + k - 1}.$$

Since $e^\alpha > 1$ for $\alpha > 0$, the probability of sending the original category is a little bit higher than sending the other category. Gaboardi and Rogers [1] constructs a private goodness-of-fit test based on a chi-square statistic evaluated on $\tilde{\mathbf{Y}}_i$'s. They demonstrate that the limiting distribution is chi-square distribution both under the null and alternative.

2.2 Two sample chi-square statistic

We extend the goodness-of-fit test by Gaboardi and Rogers [1] into two-sample testing by privatizing the raw samples $\mathbf{Z}_i \stackrel{iid}{\sim} \text{multi}(n_2, \mathbf{p}_Z)$ into $\tilde{\mathbf{Z}}_j := \mathcal{M}_{\text{GenRR}}(\mathbf{Z}_j)$. Under the null, $\mathcal{M}_{\text{GenRR}}(\mathbf{Y}_i)$ and $\mathcal{M}_{\text{GenRR}}(\mathbf{Z}_j)$ follow multinomial distributions with the same probability vector. Therefore, the usual two-sample chi-square test statistic

$$T_\chi := \sum_{\ell=1}^k \frac{(n_2 \sum_{i=1}^{n_1} \tilde{\mathbf{Y}}_i(\ell) - n_1 \sum_{j=1}^{n_2} \tilde{\mathbf{Z}}_j(\ell))^2}{n_1 n_2 (n_1 + n_2) \sum_{j=1}^{n_2} (\tilde{\mathbf{Y}}_j(\ell) + \tilde{\mathbf{Z}}_j(\ell))}$$

converges to a chi-square distribution with degree of freedom $k - 1$ and yields a valid test with size γ . This test statistic is from Van der Vaart's book Asymptotic Statistics, pp. 253.

3 Bit flip privatization and related test statisitc

3.1 Bit flip privatization

3.2 test statistic

3.2.1 Review of one-sample statistic

We first review how [1] builds goodness-of-fit chi-square statistic for histogram of bit-flipped observations. We start with applying CLT to the bit-flipped obsevation.

Lemma 3.1 (Applying CLT to bit-flipped observations, Lemma 5.7 of [1]). *When $Y_i \stackrel{iid}{\sim} \text{multinomial}(\mathbf{p}, 1)$, denote the histogram of flipped observations as*

$$\tilde{\mathbf{H}} := \sum_{i=1}^{n_1} \mathcal{M}_{bit}(Y_i). \quad (1)$$

The mean vector and covariance matrices are computed as follows:

$$\tilde{\mathbf{p}} := \mathbb{E}(\mathcal{M}_{bit}(Y_1)) = \frac{(\exp(\alpha/2) - 1)\mathbf{p} + \mathbf{1}}{\exp(\alpha/2) + 1}, \text{ and} \quad (2)$$

$$\Sigma_{\mathbf{p}} := \text{Var}(\mathcal{M}_{bit}(Y_1)) = \left(\frac{\exp(\alpha/2) - 1}{\exp(\alpha/2) + 1} \right)^2 (\text{diag}(\mathbf{p}) - \mathbf{p}\mathbf{p}^\top) + \frac{\exp(\alpha/2)}{(\exp(\alpha/2) + 1)^2} I_d, \quad (3)$$

For any $\alpha > 0$ and $\mathbf{p} > 0$, $\Sigma_{\mathbf{p}}$ is full-rank and one of its eigenvector is one-vector. By the CLT for i.i.d random vectors, we get the following asymptotic distribution:

$$\sqrt{n}(\tilde{\mathbf{H}}/n - \tilde{\mathbf{p}}) \xrightarrow{d} N(0, \Sigma_{\mathbf{p}}) \quad (4)$$

In non-private chi-square test, we apply CLT and multiply by $\text{diag}(\mathbf{p})^{-1/2}$ to turn the covariance matrix on the RHS into a projector matrix. Here in the private setting, we also need a scaling matrix to turn the covariance matrix into a projector matrix. Gaboardi and Rogers [1] proposes $\tilde{\mathbf{p}}^{-1/2}\Pi$, where $\Pi := I_k - \frac{1}{k}\mathbf{1}\mathbf{1}^\top$. The properties of Π are as follows:

1. It is symmetric idempotent (a projector matrix).
2. Its null space is $\text{span}\{\mathbf{1}\}$, so when multiplied to a symmetric matrix, it deletes an eigenvector $\mathbf{1}$.

$$\begin{aligned} \Pi x = 0 &\iff x = (1/k)\mathbf{1}\mathbf{1}^\top x \\ &\iff x = (1/k)\mathbf{1}(\mathbf{1}^\top x) = ((\mathbf{1}^\top x)/k)\mathbf{1} = c\mathbf{1} \end{aligned}$$

By multiplying $\tilde{\mathbf{p}}^{-1/2}\Pi$ to the LHS vector of (4), we get

$$\sqrt{n}\tilde{\mathbf{p}}^{-1/2}\Pi(\tilde{\mathbf{H}}/n - \tilde{\mathbf{p}}) \xrightarrow{d} N(0, \tilde{\mathbf{p}}^{-1/2}\Pi\Sigma_{\mathbf{p}}\Pi\tilde{\mathbf{p}}^{-1/2}), \quad (5)$$

where $\tilde{\mathbf{p}}^{-1/2}\Pi\Sigma_{\mathbf{p}}\Pi\tilde{\mathbf{p}}^{-1/2}$ is an identity matrix except one diagonal entry is zero. Therefore, the covariance matrix is idempotent and rank $k - 1$. Now we invoke the following classical theorem to derive an asymptotic chi-square distribution with degree of freedom $k - 1$.

Theorem 3.1 (Ferguson (1996)). *If $\mathbf{X} \sim N(\boldsymbol{\mu}, \Sigma)$ and Σ is a projection matrix of rank ν and $\Sigma\boldsymbol{\mu} = \boldsymbol{\mu}$ then $\mathbf{X}^\top \mathbf{X} \sim \chi_\nu^2(\boldsymbol{\mu}^\top \boldsymbol{\mu})$.*

We can extend this lemma to two-sample setting. Suppose $Y_i \stackrel{iid}{\sim} \text{multinomial}(\mathbf{p}_1, 1)$ and $Z_j \stackrel{iid}{\sim} \text{multinomial}(\mathbf{p}_2, 1)$. We follow Lemma 3.1 to denote $\tilde{\mathbf{p}}_Y = \mathbb{E}(\mathcal{M}_{bit}(Y_1))$, $\Sigma_{\mathbf{p}_Y} := \text{Var}(\mathcal{M}_{bit}(Y_1))$ and $\tilde{\mathbf{p}}_Z = \mathbb{E}(\mathcal{M}_{bit}(Z_1))$, $\Sigma_{\mathbf{p}_Z} := \text{Var}(\mathcal{M}_{bit}(Z_1))$. Denote $\tilde{Y}_i := \mathcal{M}_{bit}(Y_i) - \tilde{\mathbf{p}}_Y$ and $\tilde{Z}_j := \mathcal{M}_{bit}(Z_j) - \tilde{\mathbf{p}}_Z$. Then denote $T_n := \sum_{i=1}^n \tilde{Y}_i - \sum_{j=1}^n \tilde{Z}_j$ and $\Sigma_n := \text{Var}(T_n) = n(\Sigma_{\mathbf{p}_Y} + \Sigma_{\mathbf{p}_Z})$.

Under the null hypothesis of $\mathbf{p}_Y = \mathbf{p}_Z = \mathbf{p}$, we have $T_n = \sum_{i=1}^n \mathcal{M}_{bit}(Y_i) - \sum_{j=1}^n \mathcal{M}_{bit}(Z_j) = \tilde{\mathbf{H}}_Y - \tilde{\mathbf{H}}_Z$ and $\Sigma_n = 2n\Sigma_{\mathbf{p}}$. So we have

$$\sqrt{n/2}(\tilde{\mathbf{H}}_Y/n - \tilde{\mathbf{H}}_Z/n) \xrightarrow{d} N(0, \Sigma_{\mathbf{p}}) \quad (6)$$

Since $\Sigma_{\mathbf{p}}$ is symmetric and one of its eigenvector is one-vector, we can diagonalize it as $\Sigma_{\mathbf{p}} = BDB^\top$, where D is a diagonal matrix and B has orthogonal columns with one of them being $k^{-1}\mathbf{1}$.

We introduce $\Pi := I_d - \frac{1}{k}\mathbf{1}\mathbf{1}^\top$. First, this is an orthogonormal projection matrix, since it is symmetric and idempotent:

$$\begin{aligned} \Pi^2 &= \left(I_d - \frac{1}{k}\mathbf{1}\mathbf{1}^\top\right) \left(I_d - \frac{1}{k}\mathbf{1}\mathbf{1}^\top\right) = I_d - \frac{1}{k}\mathbf{1}\mathbf{1}^\top - \frac{1}{k}\mathbf{1}\mathbf{1}^\top + \frac{1}{k^2}\mathbf{1}\mathbf{1}^\top\mathbf{1}\mathbf{1}^\top \\ &= I_d - 2\frac{1}{k}\mathbf{1}\mathbf{1}^\top + \frac{1}{k^2}\mathbf{1}(\mathbf{1}^\top\mathbf{1})\mathbf{1}^\top \\ &= I_d - 2\frac{1}{k}\mathbf{1}\mathbf{1}^\top + \frac{1}{k^2}\mathbf{1}(k\mathbf{1}^\top) \\ &= I_d - \frac{1}{k}\mathbf{1}\mathbf{1}^\top \\ &= \Pi. \end{aligned}$$

$\mathbf{1}\mathbf{1}^\top$ Since Π is symmetric, its column space is the orthogonal complement of $\text{span}\{\mathbf{1}\}$. So multiplying by Π means under the null, it suffices to use the CLT for i.i.d. random vectors, but under the alternative, we would need to use Lindeburg or Lyapunov.

$$\frac{\tilde{\mathbf{H}}_1}{n_1} - \frac{\tilde{\mathbf{H}}_2}{n_2}$$

References

- [1] Gaboardi, M. and Rogers, R. (2018). Local private hypothesis testing: Chi-square tests. *Proceedings of the 35th International Conference on Machine Learning*, 80:1626–1635.